



Applied Parallel Computing
parallel-computing.pro

Memory Hierarchy

Pavel Yakimov



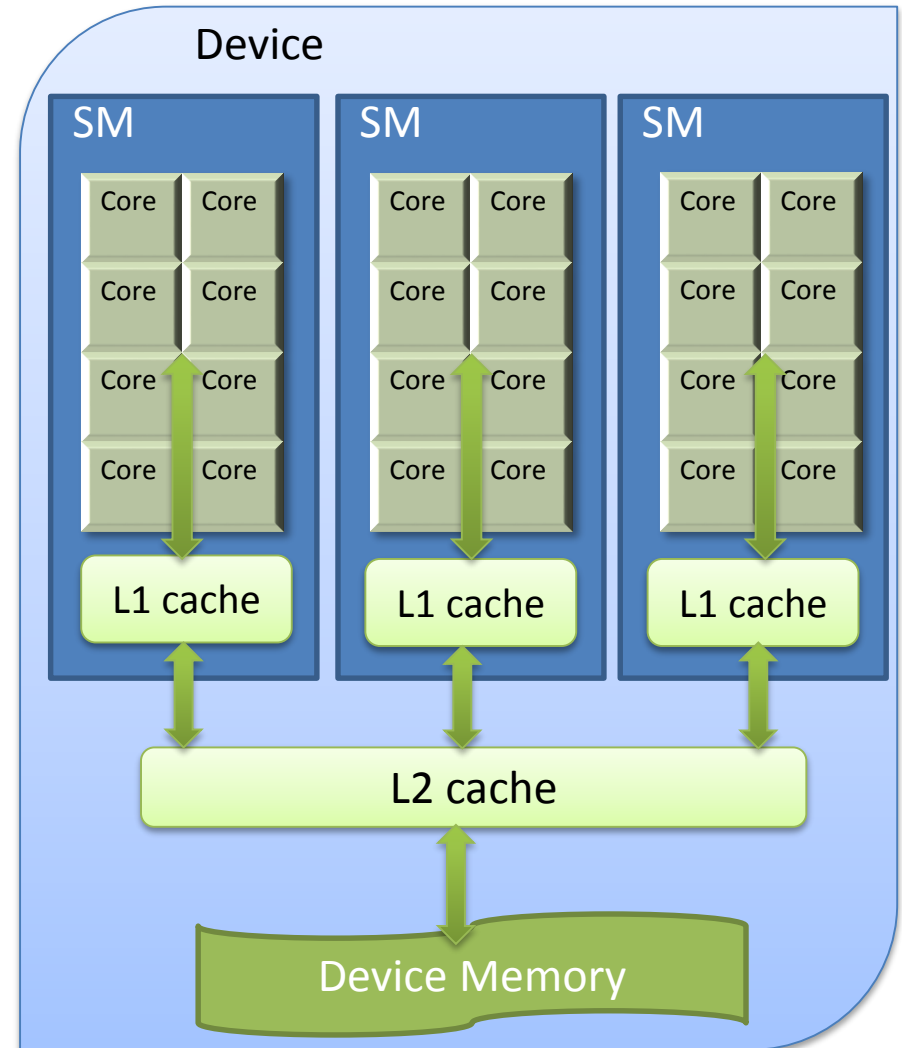
Applied Parallel Computing
parallel-computing.pro

Global Memory




Global Memory

- Located in **DRAM GPU**
- Up to 16Gb
 - Can be checked by `totalGlobalMem`
- **Cacheable**, uses caches L1 and L2:
 - L1 – located in each SM
maximum size - 64KB
minimal size - 0KB
 - L2 – on device
maximum size 6 MB
Device parameter `l2CacheSize`





L1 Management

- ④ Caching to L1 can be switched off **at compilation**
 - `nvcc -Xptxas -dlcm=ca`
enabled L1 caching (by default)
 - `nvcc -Xptxas -dlcm=cg`
disabled L1 caching, in binary code all global memory accesses are translated to instructions that don't use L1 while executing 
- ④ The differences are in binary code, not in execution mode





Transactions

- ④ Access to global memory:
 - With enabled cache L1 – transactions of 128 bytes
 - ✓ Cache line L1 – 128 bytes
 - With disabled cache L1 – transactions of 32 bytes
 - ✓ Cache line L2 – 32 bytes
- ④ Transactions are aligned by size(naturally aligned)
- ④ Memory access instruction is performed simultaneously for all threads in a warp (SIMT)
 - Threads can access to different elements
 - Performing as many transactions, as needed to cover memory requests of all threads in a warp




Access Patterns: L1 enabled





Access Patterns: L1 enabled

Warp requests




Requested 128 bytes, but the address is unaligned

2 transactions - 256 bytes



Warp requests



Requested 128 bytes, but requests are scattered within 3 cache lines -


3 transactions – 384 bytes





Access Patterns: L1 disabled

Warp requests



Requests 128 bytes, but the address is unaligned. Fits to 4 cache lines in L2
4 transactions of 32 bytes each – 128 bytes

Warp requests



Requests 128 bytes, but requests are scattered within 5 cache lines
5 transactions of 32 bytes each – 160 bytes




Matrices in Global Memory

- ⌚ A matrix is stored linearly, row-by-row
- ⌚ Let the length of a row – 480 bytes (120 float)
 - access – `matrix[idy*120 + idx]`

Address of each row beginning, except the first one, is not aligned to 128 bytes–

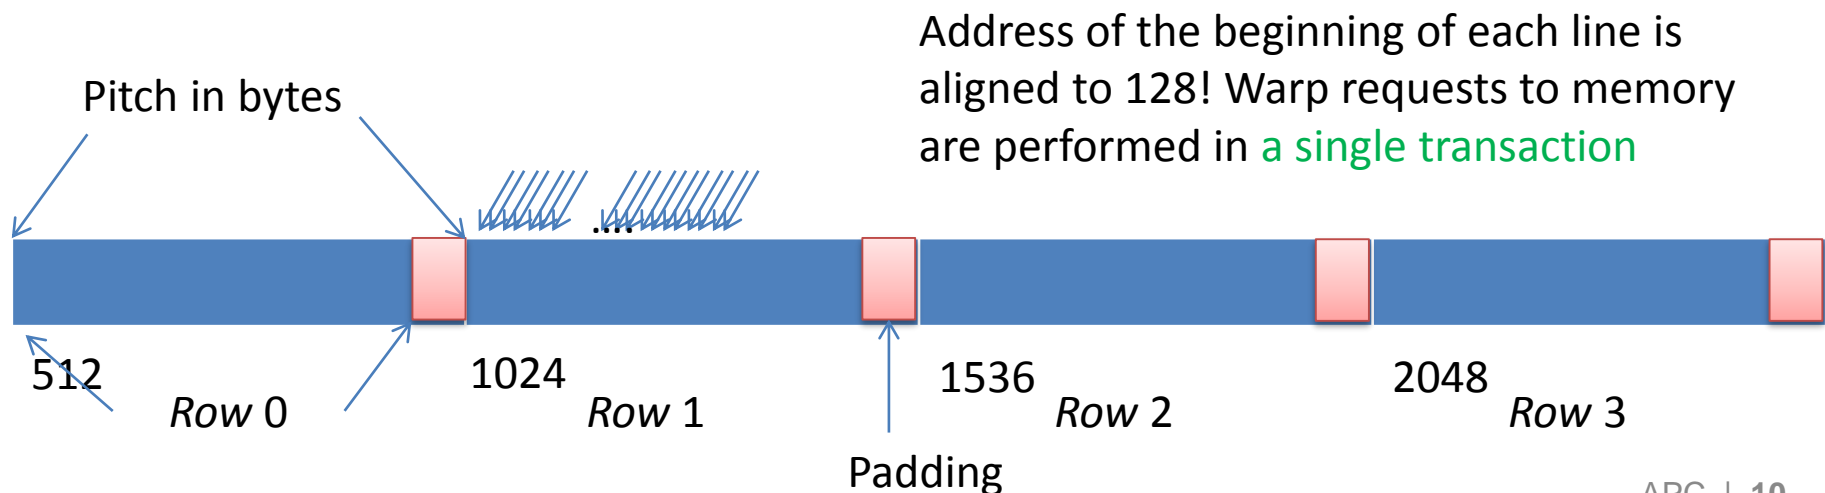
2 transactions





Matrices in Global Memory

- Supplement each row to a multiple of 128 bytes – in this example, $480 + 32 = 512$,
32 bytes is **pitch** – actual width in bytes
- These bytes can not be used, i.e. $32/512=6\%$ extra memory will be allocated (but for large matrices, this share will be significantly less)
- However, each row will be aligned to 128 bytes
 - Request `matrix[idy*128+ idx]`





Memory Allocation with Padding

- ④ `cudaError_t cudaMallocPitch (void ** devPtr, size_t * pitch, size_t width, size_t height)`
 - `width` – logical matrix width in `bytes`
 - Allocates not less than `width * height bytes`, can add some padding to the end of a row, in order to align the beginning of rows
 - Stores the pointer to memory (`*devPtr`)
 - Stores actual width of rows to (`*pitch`)
- ④ Matrix element address (Row, Column), allocated with `cudaMallocPitch`:

`T* pElement = (T*) ((char*) devPtr + Row * pitch) + Column;`





Copying to Matrix with Padding

⑧ `cudaError_t cudaMemcpy2D (void* dst, size_t dpitch, const
void* src, size_t spitch, size_t width,
size_t height, cudaMemcpyKind kind)`

- `dst` - a pointer to the matrix you want to copy ,
`dpitch` – **actual** row width in bytes
- `src` - a pointer to the matrix from which you want to copy,
`spitch` – **actual** row width in bytes
- `width` – width of matrix transfer (columns in bytes)
- `height` – Height of matrix transfer (rows)
- `kind` – type of transfer (similar to `cudaMemcpy`)

⑧ This function `width` bytes from the beginning of each row of source matrix. Total transfer size is `width*height` bytes, herewith

- Address of row with index `Row` is calculated with actual width:


`(char*) src + Row* spitch` – *source matrix*

`(char*) dst + Row* dpitch` – *destination matrix*



How to Reference to Matrix by Columns?

- Matrix is allocated by rows, and the memory requests are by columns



If the matrix has a size greater than 128 bytes, then these requests would never fit to a single transaction!



Transpose!


- **Solution – store transposed matrix!**
 - This leads into actual sequential addressing to rows in memory just like they are columns
 - The memory for transposed matrix should also be allocated with `cudaMallocPitch`



Array of Structs?

```
struct example {  
    int a;  
    int b;  
    int c;  
}  
  
__global__ void kernel(example * arrayOfExamples) {  
    int idx = threadIdx.x + blockIdx.x * blockDim.x;  
    arrayOfExamples[idx].c =  
        arrayOfExamples[idx].b + arrayOfExamples[idx].a;  
}
```

Warp requests




Warp addressing to memory will
execute in 3 transactions - **256**
extra bytes



Struct of Arrays!

```
struct example {  
    int *a;  
    int *b;  
    int *c;  
}  
  
__global__ void kernel(example arrayOfExamples) {  
    int idx = threadIdx.x + blockIdx.x * blockDim.x;  
    arrayOfExamples.c[idx] =  
        arrayOfExamples.b[idx] + arrayOfExamples.a[idx];  
}
```



Warp memory requests fit to
one transaction



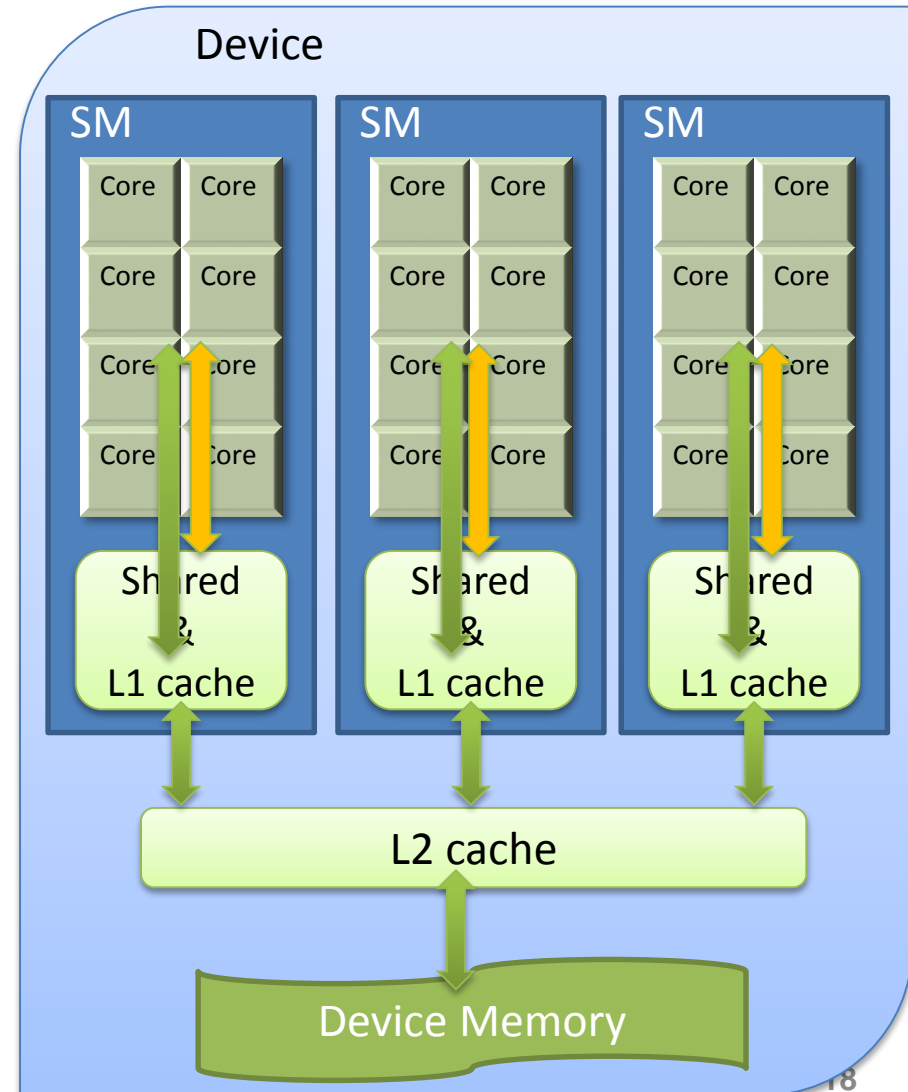
Applied Parallel Computing
parallel-computing.pro

Shared Memory



Shared Memory

- Located in the same device in SM as L1
- Jointly used by all threads in a block
- If a multiprocessor executes several blocks - memory is divided equally between them
- Each block gets its limited address space of shared memory
- Configurations:
 - 16KB shared memory, 48KB L1
 - 32KB shared memory, 32KB L1
 - 48KB shared memory, 16KB L1 – by default





Shared Memory Allocation

Statically:

- In GPU code declare a static array or a variable with qualifier `__shared__`

```
#define SIZE 1024
__shared__ int array[SIZE]; //array
__shared__ float varSharedMem; //variable
```

Dynamically:

- In GPU code declare a pointer to array in shared memory :

```
extern __shared__ int array[];
```

- In host code specify how much shared memory in bytes should be additionally allocated per each block. The third parameter of kernel launch passes this value

```
kernel<<<gridDim, blockDim, SIZE >>>(params)
```



Features of Use

- ④ In terms of programming, variables with qualifier **__shared__** :
 - Can be declared in the global scope or within functions
 - When declared in a function perform as static, i.e. one instance exists for all function calls
 - Individual for each block and attached to their personal space of shared memory
 - each block of threads sees **'his'** value
 - Exist only for the lifetime block
 - not available from host or from other blocks
 - **Can not be initialized when declaring**



Features of Use

- Shared memory is statically allocated in kernel only if it uses any `__shared__` variables, declared without `[]` in
 - the global scope
 - in kernel
 - In functions, called from kernel,
- Kernel can be simultaneously provided by both dynamically and statically allocated shared memory
- All variables `extern __shared__ type var[]` point to the beginning dynamic shared memory allocated to block



Synchronization

- Consider the example of kernel running on a linear one-dimensional grid :

```
__global__ void kernel() {  
    __shared__ int shmem[BLOCK_SIZE];  
    shmem[threadIdx.x] = __sinf(threadIdx.x);  
    int a = shmem[(threadIdx.x + 1) % BLOCK_SIZE];  
    ...  
}
```

- Each thread
 - Calculates `__sinf` from its index and stores it to the corresponding element of the output array
 - Reads an element written by neighbor thread



Synchronization

- Consider the example of kernel running on a linear one-dimensional grid:

```
__global__ void kernel() {  
    __shared__ int shmem[BLOCK_SIZE];  
    shmem[threadIdx.x] = __sinf(threadIdx.x);  
    int a = shmem[(threadIdx.x + 1) % BLOCK_SIZE];  
    ...  
}
```

- Warp execute in unpredictable order

- It may happen that a thread has not yet written a value, and the neighboring thread already tries to read it!
- *read-after-write, write-after-read, write-after-write* conflicts



Synchronization

For explicit synchronization some built-in functions are provided :

- `void __syncthreads () ;`

When you call this function, thread waits until:

- ✓ all threads in a block reach this point
- ✓ the results of all currently initiated operations with global / shared memory will be visible to all threads of the block

- `__syncthreads()` can be invoked in the branches of 'if' statement only if the result of its condition is the same for all threads in a block,
- if not then execution may deadlock or become unpredictable



Synchronization

```
__global__ void kernel() {  
    __shared__ int shmem[BLOCK_SIZE];  
    shmem[threadIdx.x] = __sinf(threadIdx.x);  
    __syncthreads();  
    int a = shmem[(threadIdx.x + 1) % BLOCK_SIZE];  
    ...  
}
```

- Each thread
 - Calculates `__sinf` from its index and stores it to the corresponding element of the output array
 - Wait until all threads in a block complete calculations
 - Reads an element written by neighbor thread



Strategy of Usage

- Shared memory can be considered as a cache, **controlled by a programmer**
 - Low latency – located on the same chip as L1, speed of memory requests is comparable to registers
 - Application explicitly allocates and uses shared memory
 - Access pattern can be arbitrary, unlike in L1
- Even if the hardware cache L1 is able to process all requests (L1 load hit $\sim 100\%$), the use of shared memory allows making full use of the equipment
 - Otherwise 16KB (or 48KB, if you forgot to set the proper mode) fast shared memory is idle



Strategy of Usage

- Typical strategy:
 - Threads in a block **collectively**:
 1. Download data from global memory to shared
 - Each thread executes part of this downloading
 2. Synchronize
 - That no thread starts reading the data being uploaded by another thread, before it finishes uploading
 3. Use the downloaded data to calculate results
 - If threads write something to shared memory, it may also need to synchronize again
 4. Write the results to global memory



Example of Shared Memory Usage

```
__global__ void kernel(int sizeofArray1, int sizeofArray2, int *devPtr, int *res)
{
    extern __shared__ int dynamicMem[]; // a pointer to dynamic shared memory
    __shared__ int staticMem[1024];    // static array in shared memory
    __shared__ int var;                // a variable in shared memory
    int *array1 = dynamicMem; // address of the first array in dynamic shared memory
    int *array2 =
        array1 + sizeofArray1; // address of the second array in dynamic shared memory

    staticMem[threadIdx.x] = devPtr[threadIdx.y]; // loading data to shared memory
    __syncthreads(); // wait until all threads finish loading

    array2[threadIdx.x] =
        2 * staticMem[(threadIdx.x - 10) % blockDim.x]; // access to an element
                                                         // written by another thread

    __syncthreads(); // wait until all threads finish writing
    res[threadIdx.x] = array2[threadIdx.x]; // write the results to global memory
}
```



Example of Shared Memory Usage

```
__global__ void kernel(int sizeOfArray1, int sizeOfArray2, int *devPtr, int *res)
{
    extern __shared__ int dynamicMem[]; // a pointer to dynamic shared memory
    __shared__ int staticMem[1024];    // static array in shared memory
    __shared__ int var;                // a variable in shared memory

    int *array1 = dynamicMem; // address of the first array in dynamic shared memory
    int *array2 =
        array1 + sizeOfArray1; // address of the second array in dynamic shared memory

    staticMem[threadIdx.x] = devPtr[threadIdx.y]; // loading data to shared memory
    __syncthreads(); // wait until all threads finish loading

    array2[threadIdx.x] =
        2 * staticMem[(threadIdx.x - 10) % blockDim.x]; // access to an element
                                                         // written by another thread

    // Actually, there is no need in synchronization here
    res[threadIdx.x] = array2[threadIdx.x]; // write the results to global memory
}
```



Example of Kernel Launching with Dynamic Shared Memory

Host code

```
int *devPtr;  
cudaMalloc(&devPtr, 1024*sizeof(int));  
kernel<<<3,1024,1024*sizeof(int)>>>(512,512, devPtr); // launch three blocks of  
//1024 threads. Dynamically allocate 4KB of shared memory per block
```

Size of dynamically allocated
memory, in bytes per block

- If the total (static + dynamic) requested shared memory size exceeds the total available memory (16KB, 32KB or 48KB), an error will occur



Writing to Shared Memory

- ④ Several threads of a warp trying to write to the same address
 - The operation will be executed only by one thread
 - Which – **unknown**
 - ✓ (Apparently - the latter thread of a warp, from those that have to write)
 - ✓ The result is unpredictable
 - At least, because the order is unpredictable and no one knows what warp will be the latest



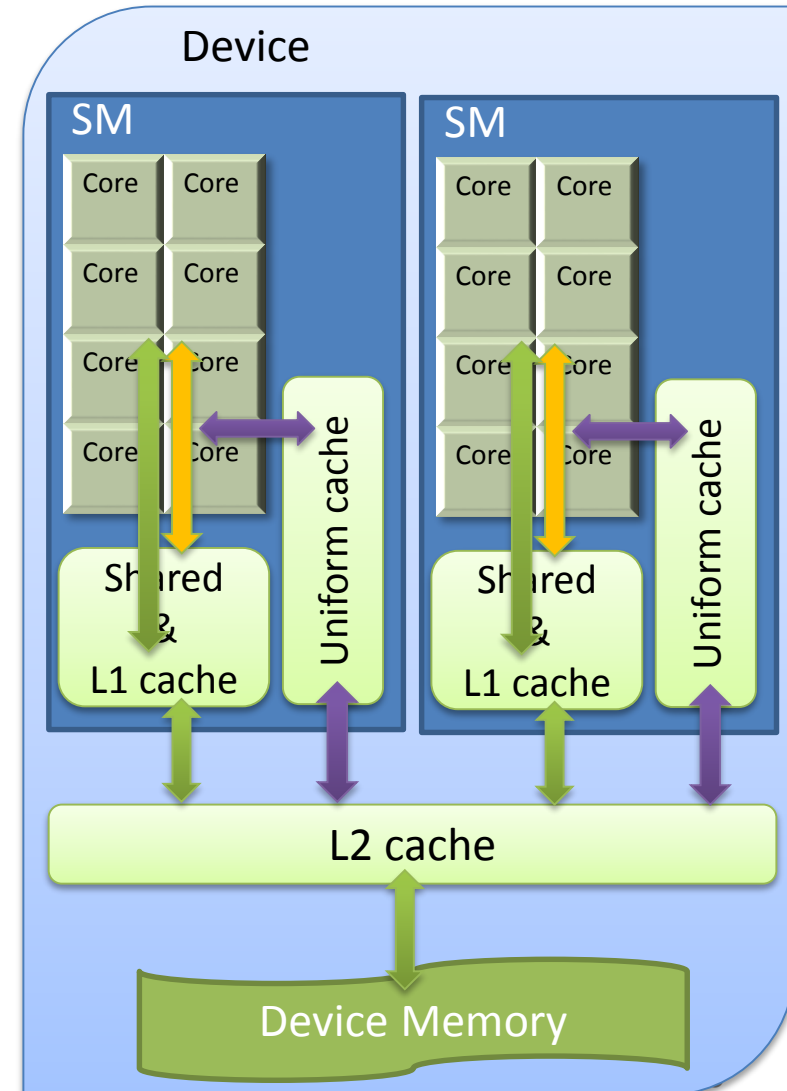
Applied Parallel Computing
parallel-computing.pro

Constant Memory



Constant Memory

- Located in **DRAM GPU**
- Size up to 64KB
 - Device parameter **totalConstMem**
- **Cacheable** in special read-only cache – Uniform Cache
 - Up to 8 KB





Declaration

- In the global scope

```
__constant__ int constMem[1024];  
__constant__ int constVar;
```

- You can additionally specify `__device__`, to indicate that the memory is allocated on the device:

```
__device__ __constant__ int constVar2;
```



Features

- Allocated at application startup, released at the end of application lifetime
- Is available for reading (**read-only!**) from any thread of any grid in a usual way:

```
__constant__ int constMem[32];  
__global__ void kernel() {  
    ...  
    int a = constMem[ threadIdx.x / 32 ];  
    ...  
}
```

- Available from host with special functions from toolkit:
`cudaGetSymbolAddress()` / `cudaGetSymbolSize()` / `cudaMemcpyToSymbol()` / `cudaMemcpyFromSymbol()`



Example

```
__constant__ float constData[256];
```

 Host:

```
float data[256];  
cudaMemcpyToSymbol(constData, data, sizeof(data));  
cudaMemcpyFromSymbol(data, constData, sizeof(data));
```



Requests to Constant Memory

- Request is performed simultaneously for all threads in a warp (SIMT)
- The initial requests is divided into as many queries as many different addresses it contains
 - Each query is performed either through a request to cache, in case of cache hit, or to global memory
 - If there were n queries, then the bandwidth is reduced by n times

```
__constant__ int constMem[32];  
__global__ void kernel() {  
    ...  
    int a = constMem[ threadIdx.x / 32 ]; // 1 request to constant  
                                         memory per warp  
    int a = constMem[ threadIdx.x ]; // 32 request to constant  
                                     memory per warp  
    ...  
}
```



Uniform Access

- In addition to processing requests to constant memory, Uniform Cache also processes the Uniform Accesses - when all threads of a single warp accesses global memory at the same address

- It is possible only when:
 - ✓ The access is read-only
 - ✓ The address doesn't depend on the thread's index (**threadIdx**)

```
while (k < 100 ) tmp += a[blockIdx.x + k++];
```

In assembler, compiler will change an ordinary instruction of load from global memory to the instruction of uniform load, which will be executed using Uniform Cache

- The second requirement guarantees that all threads of a warp request to the same address
- To help the compiler with the first requirement, we can add the qualifier **const** to the pointers



Passing Parameters and Grid to Kernels

- Parameters are passed to the kernel through the constant memory
 - Parameters are passed in a single copy for all threads of grid
 - This is acceptable, because,
 - ✓ basically, threads of a warp are requesting the same parameter -> **Uniform Access**
 - ✓ after the first warp, the parameters **will already be in the cache**
- The total size the passed parameters must be **no larger than 4 KB**
- Grid size (**gridDim**, **blockDim**) is also passed through the constant memory :
 - Thread receives **threadIdx** and **blockIdx** from special registers
 - gridDim**, **blockDim** are **read from the constant memory** at the very beginning of execution (Uniform)



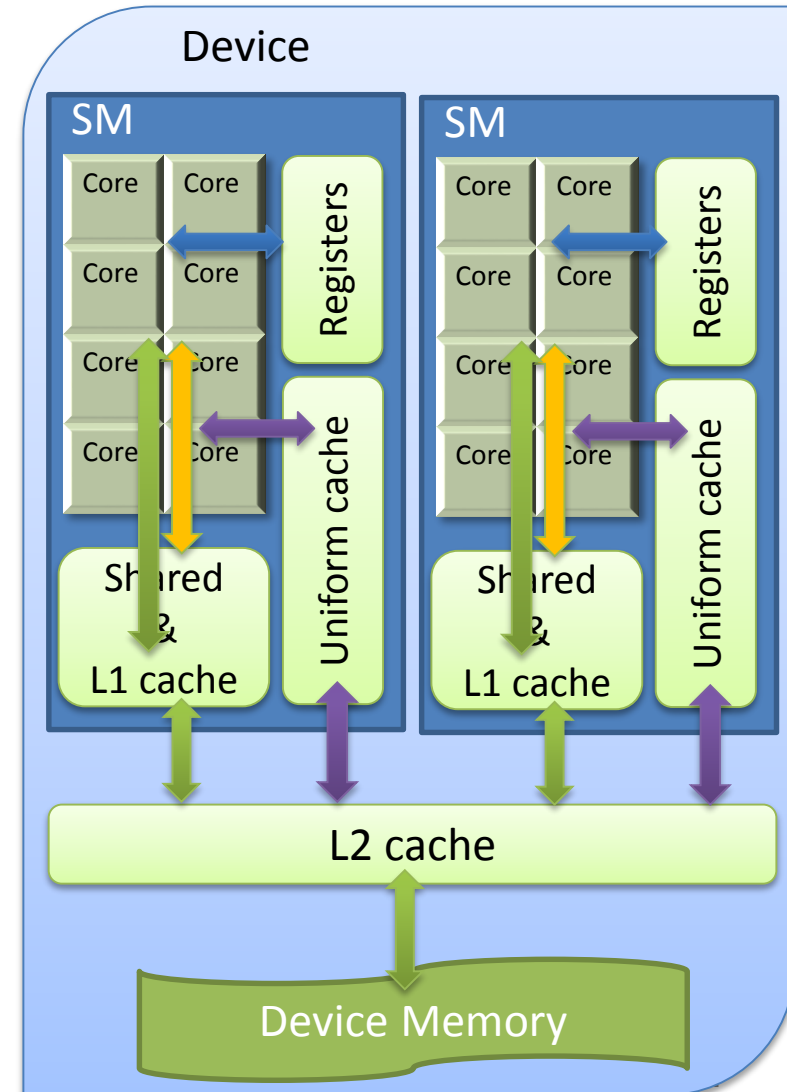
Applied Parallel Computing
parallel-computing.pro

Registers and Local memory



Registers

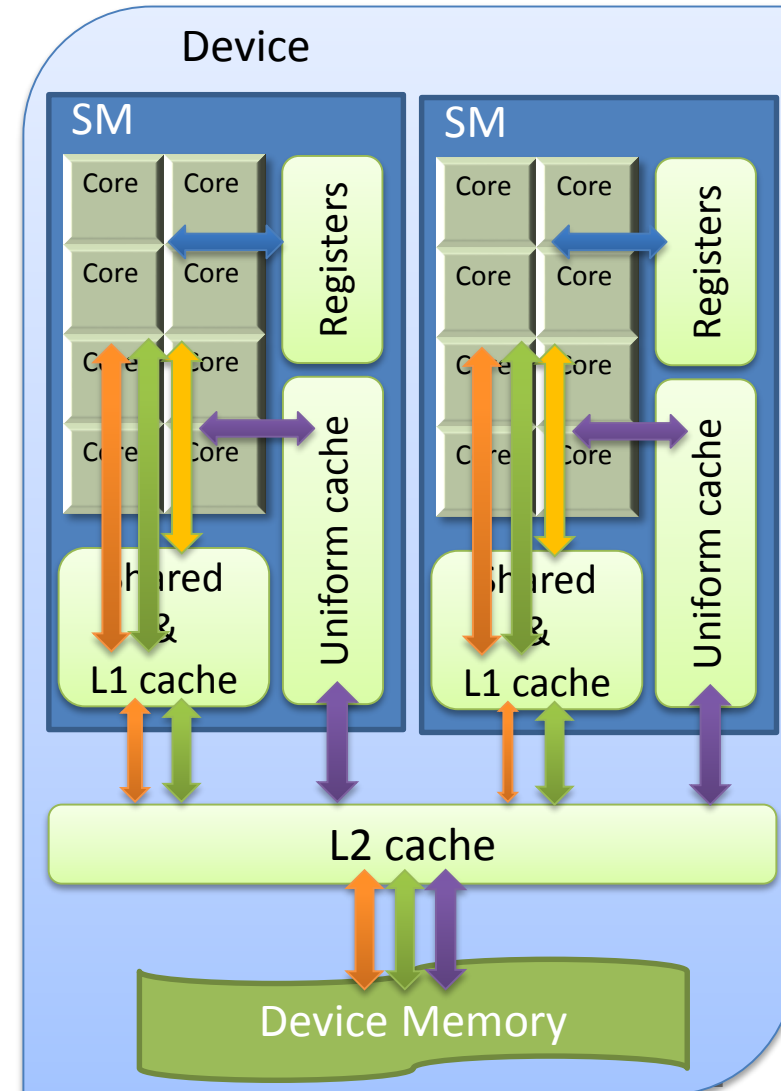
- Located on SM
- The fastest memory access
- Each multiprocessor contains 64K 32-bit registers
 - 256 KB of registers
 - Device parameter - **regsPerBlock**
- One thread can have up to 63 registers
- Distributed **during compilation**
- Each thread is an exclusive user of its registers for the duration of the kernel execution. Access to the registers of other threads is restricted





Local Memory

- Located in **DRAM**
- Access is made by the same rules as requests to the global memory
 - Caching to L1
 - Transactions
- Unavailable explicitly
- Has simplified addressing scheme
 - Optimized to minimize the number of transactions





When is local memory used?

- Typically, the compiler places into registers all local variables
- But there are exceptions, which are placed into the local memory
 - Arrays, for which you can not always determine which element in what period of time is being accessed (not constant indexes)
 - Large arrays or structures that would use too many registers
 - Any variable, if the limit of 63 registers per thread is exceeded (register spilling)
- Some built-in math functions can use the local memory
- Local memory is used to pass a part of the operands for function call
 - Stack frame for recursive calls is modeled in the local memory



nvcc -Xptxas -v

- Displays the number of registers, constant memory, local memory and static shared memory used by the kernel:

```
pyakimov@localhost:~/programming/testMod$ nvcc -arch=sm_20 -Xptxas -v test.cu
ptxas info    : 0 bytes gmem, 8 bytes cmem[2]
ptxas info    : Compiling entry function '_Z13matmul_kernelv' for 'sm_20'
ptxas info    : Function properties for _Z13matmul_kernelv
    8 bytes stack frame, 0 bytes spill stores, 0 bytes spill loads
ptxas info    : Used 8 registers, 4 bytes smem, 32 bytes cmem[0]
```



Thank You!



L1 Cache Configuration

- Cache supports 3 configuration modes: 48KB, 32KB , and 16KB
- Switching modes:

- `cudaError_t cudaDeviceSetCacheConfig(cudaFuncCache cacheConfig)`

`cacheConfig` resets cache mode on the whole device

✓ Possible modes:

- `cudaFuncCachePreferNone` – without preferences (by default). Sets the last used configuration . The initial configuration – 16KB L1
- `cudaFuncCachePreferShared`: 16KB L1
- `cudaFuncCachePreferEqual`: 32KB L1
- `cudaFuncCachePreferL1`: 48KB L1

- `cudaError_t cudaFuncSetCacheConfig (const void* func, cudaFuncCache cacheConfig)`

Switches cache to `cacheConfig` configuration for function `func`

✓ By default - `cudaFuncCachePreferNone`



L1 Cache Configuration

- It is just 'preferences'
 - Runtime may not provide the preset configuration mode depending on kernel and device parameters
- If a kernel doesn't use shared memory, then it is favorable to switch to `cudaFuncCachePreferL1` mode
- In general, it is worth checking the performance of all modes:
 - `(-dlcm=ca, -clcm=cg)x(16KB, 32KB, 48KB)`

